

Introduction:

The data set I chose is Movie dataset

Here are some notes and comments about this datasets : This data set contains information about 10,000 movies collected from The Movie Database (TMDb), including user ratings and revenue.

Certain columns, like 'cast' and 'genres', contain multiple values separated by pipe (|) characters.

There are some odd characters in the 'cast' column. Don't worry about cleaning them. You can leave them as is.

The final two columns ending with "_adj" show the budget and revenue of the associated movie in terms of 2010 dollars, accounting for inflation over time.

The questions about this dataset:

1. Does higher budget mean higher popularity ? Is there a coefficient relationship ?
2. Will the runtime affect the vote count and popularity?
3. Higher popularity means higher profits ?
4. What Features are Associate with Top 10 Revenue Movies ?
5. Which genres are most popular from year to year?

▼ Data Wrangling:

Get familiar with the data types, data structure. I did delete the duplicates and unuseful columns like imdb_id,homepage etc.

When handling the missing data. I use two ways: for all the missing data with data type object, i fill the null with string "missing". For budget, datatype integer,I fill 0 with np.NaN.

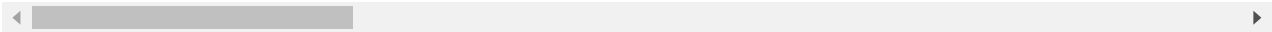
```
from google.colab import drive
drive.mount('/content/drive')
```

```
# Import all the libraries which will be used later
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
%matplotlib inline
```

```
# load in data and print out the head
df=pd.read_csv('/content/drive/MyDrive/Colab Notebooks/tmdb-movies.csv')
df.head()
```

	id	imdb_id	popularity	budget	revenue	original_title	
0	135397	tt0369610	32.985763	150000000	1513528810	Jurassic World	Chris Pratt Howe k
1	76341	tt1392190	28.419936	150000000	378436354	Mad Max: Fury Road	Hardy Thei By
2	262500	tt2908446	13.112507	110000000	295238201	Insurgent	Wood Jar Winsle
3	140607	tt2488496	11.173104	200000000	2068178225	Star Wars: The Force Awakens	F Harr Fisher A
4	168259	tt2820852	9.335014	190000000	1506249360	Furious 7	Vin Die Walk Statham

5 rows × 21 columns



```
# check the rows and columns of this dataset
df.shape

(10866, 21)
```

```
# check datatypes to see if there are some wrongly categorized types
df.dtypes
```

id	int64
imdb_id	object
popularity	float64

```
budget          int64
revenue         int64
original_title  object
cast            object
homepage        object
director        object
tagline         object
keywords        object
overview        object
runtime         int64
genres          object
production_companies object
release_date    object
vote_count      int64
vote_average    float64
release_year    int64
budget_adj      float64
revenue_adj     float64
dtype: object
```

```
# check each columns number of unique values
df.nunique()
```

```
id              10865
imdb_id         10855
popularity      10814
budget          557
revenue         4702
original_title  10571
cast            10719
homepage        2896
director        5067
tagline         7997
keywords        8804
overview        10847
runtime         247
genres          2039
production_companies 7445
release_date    5909
vote_count      1289
vote_average     72
release_year     56
budget_adj      2614
revenue_adj     4840
dtype: int64
```

```
# statistic values for this data
df.describe()
```

	id	popularity	budget	revenue	runtime	vote
count	10866.000000	10866.000000	1.086600e+04	1.086600e+04	10866.000000	10866.
mean	66064.177434	0.646441	1.462570e+07	3.982332e+07	102.070863	217.
std	92130.136561	1.000185	3.091321e+07	1.170035e+08	31.381405	575.
min	5.000000	0.000065	0.000000e+00	0.000000e+00	0.000000	10.
25%	10596.250000	0.207583	0.000000e+00	0.000000e+00	90.000000	17.
50%	20660.000000	0.382856	0.000000e+00	0.000000e+00	90.000000	20.

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10866 entries, 0 to 10865
Data columns (total 21 columns):
#   Column                               Non-Null Count  Dtype
---  -
0   id                                    10866 non-null  int64
1   imdb_id                             10856 non-null  object
2   popularity                           10866 non-null  float64
3   budget                              10866 non-null  int64
4   revenue                              10866 non-null  int64
5   original_title                       10866 non-null  object
6   cast                                 10790 non-null  object
7   homepage                             2936 non-null  object
8   director                             10822 non-null  object
9   tagline                              8042 non-null  object
10  keywords                             9373 non-null  object
11  overview                             10862 non-null  object
12  runtime                              10866 non-null  int64
13  genres                               10843 non-null  object
14  production_companies                 9836 non-null  object
15  release_date                         10866 non-null  object
16  vote_count                           10866 non-null  int64
17  vote_average                         10866 non-null  float64
18  release_year                         10866 non-null  int64
19  budget_adj                           10866 non-null  float64
20  revenue_adj                          10866 non-null  float64
dtypes: float64(4), int64(6), object(11)
memory usage: 1.7+ MB
```

```
df.isnull().sum()
```

```
id                0
imdb_id           10
popularity        0
budget            0
revenue           0
original_title    0
cast              76
homepage          7930
director          44
tagline           2824
keywords          1493
overview          4
runtime           0
```

```

genres                23
production_companies  1030
release_date          0
vote_count            0
vote_average          0
release_year          0
budget_adj            0
revenue_adj           0
dtype: int64

```

```
# drop unuseful columns
```

```
df.drop(['id','imdb_id', 'homepage','overview'],axis=1,inplace=True) # do not forget
```

```
# Ways to handle missing data
```

```
# For all missing data with object as datatype , I fill in with string "missing"
```

```
df['cast'].fillna('missing',inplace=True )
```

```
df['director'].fillna('missing',inplace=True)
```

```
df['tagline'].fillna('missing',inplace=True)
```

```
df['keywords'].fillna('missing',inplace=True)
```

```
df['genres'].fillna('missing',inplace=True)
```

```
df['production_companies'].fillna('missing',inplace=True)
```

```
df['budget'] = df['budget'].replace(0, np.NaN)
```

```
# although there is no null in budget, but we would find there is a problem when we
```

```
# Will deal with all the 0 value in budget later.
```

```
# confirm the data
```

```
df.isnull().sum()
```

```

popularity            0
budget                5696
revenue               0
original_title        0
cast                  0
director              0
tagline               0
keywords              0
runtime               0
genres                0
production_companies  0
release_date          0
vote_count            0
vote_average          0
release_year          0
budget_adj            0
revenue_adj           0
dtype: int64

```

```
# check if there are some duplicates
```

```
df.duplicated().sum()
```

```
1
```

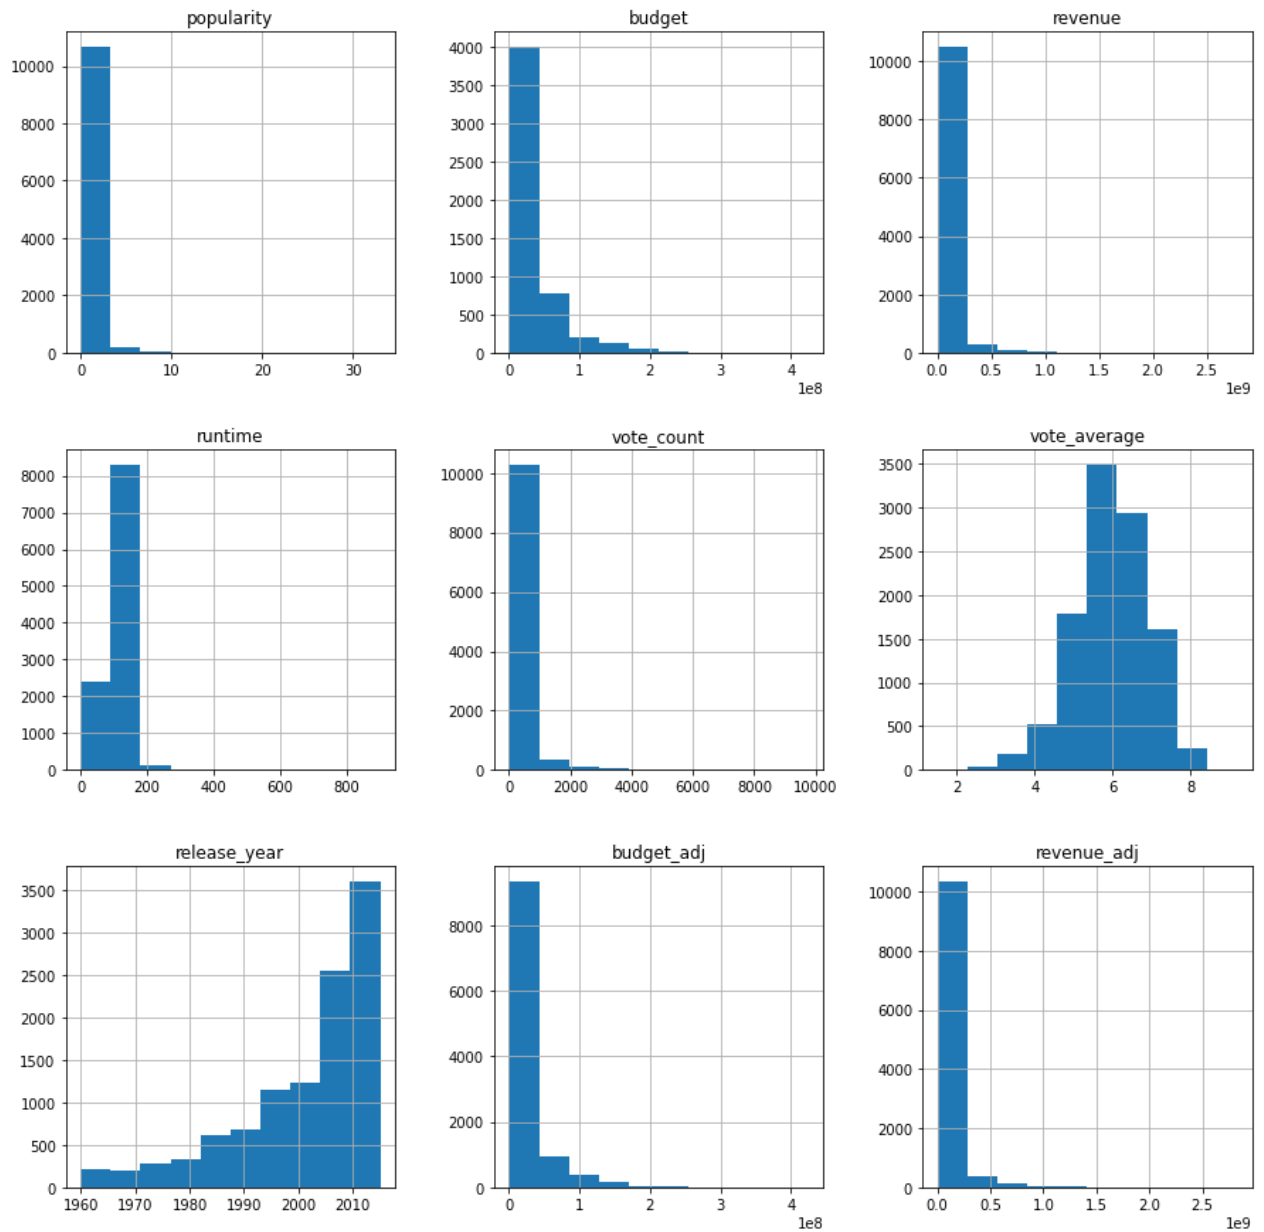
```
# drop the duplicates
```

```
df.drop_duplicates(inplace=True) # do not forget inplace = True
```

```
# confirm again  
df.duplicated().sum()
```

0

```
# visulize each variables  
df.hist(figsize=(15,15));
```

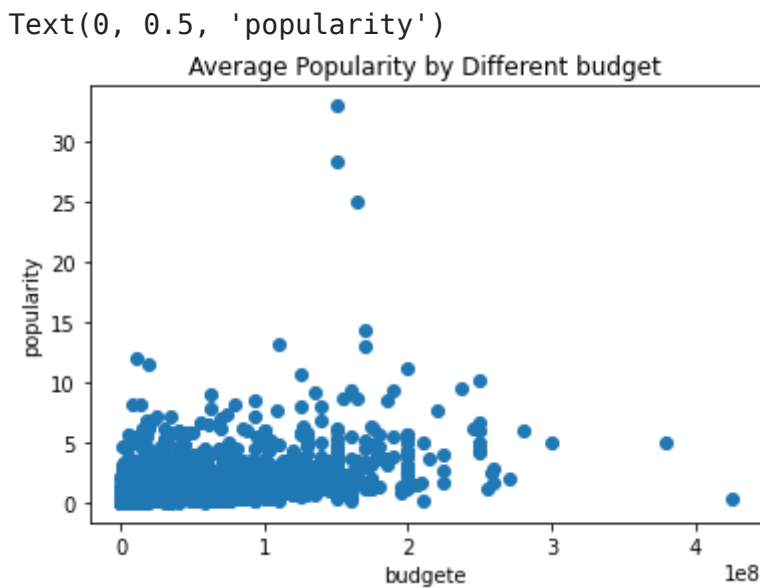


▼ Exploration with Visuals and Conclusions

Question 1. Does higher budget mean higher popularity ? Is there a coefficient relationship ?

```
# plot the relation between budget and popularity
x = df['budget']
y = df['popularity']

plt.scatter(x,y)
plt.title('Average Popularity by Different budget',fontsize=12)
plt.xlabel('budgete',fontsize=10)
plt.ylabel('popularity',fontsize=10)
```



We can not see very strong relationship between the budget and the popularity from above plot. Let's try to compare the data in another way: create two groups based on median value of budget

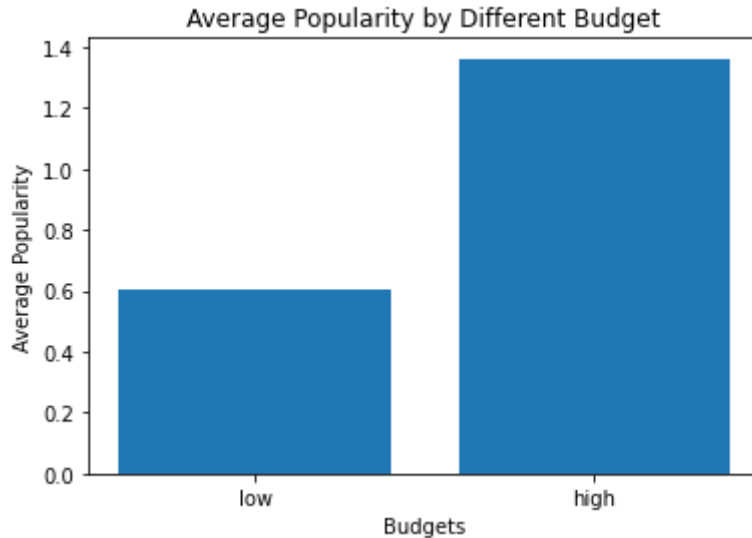
```
# based on median budget value to divide the budget into two groups : low and high
m = df['budget'].median()
low_budg = df.query('budget < {}'.format(m))
high_budg = df.query('budget >= {}'.format(m))
```

```
# check low budget and high budget mean values respectively
mean_popularity_of_low_budget = low_budg['popularity'].mean()
mean_popularity_of_high_budget = high_budg['popularity'].mean()
```

```
# create a bar chart with the values we get above
locations = [1,2]
heights = [mean_popularity_of_low_budget , mean_popularity_of_high_budget]
labels=['low','high']
plt.bar(locations, heights, tick_label = labels)
plt.title('Average Popularity by Different Budget')
```

```
plt.xlabel('Budgets')
plt.ylabel('Average Popularity')
```

```
Text(0, 0.5, 'Average Popularity')
```



```
increase_percentage = (mean_popularity_of_high_budget - mean_popularity_of_low_budget) / mean_popularity_of_low_budget * 100
increase_percentage
```

```
55.50933772947093
```

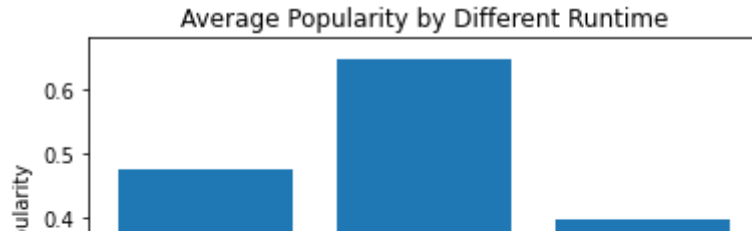
```
# here I will create 3 groups with query(). <60 min: short , 60 min <= 120 min: medium, >120 min: long
short = df.query('runtime < {}'.format(100))
medium = df.query('runtime < {}'.format(200))
long = df.query('runtime > {}'.format(200))
```

```
# check mean popularity of different movie lengths
mean_popularity_of_short = short['popularity'].mean()
mean_popularity_of_medium = medium['popularity'].mean()
mean_popularity_of_long = long['popularity'].mean()
```

```
locations = [1,2,3]
heights = [mean_popularity_of_short, mean_popularity_of_medium, mean_popularity_of_long]
labels=['low','medium','high']
plt.bar(locations, heights, tick_label = labels)
plt.title('Average Popularity by Different Runtime')
plt.xlabel('Runtime')
plt.ylabel('Average Popularity')
```



```
Text(0, 0.5, 'Average Popularity')
```



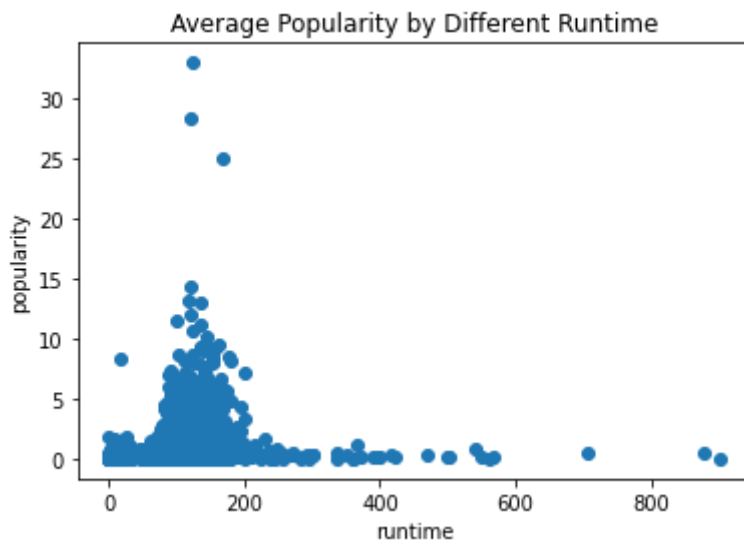
The movies should not be too long or too short. Medium length is better to gain higher popularity. But above bar chart is hard to tell the best length of runtime. Scatter plot may be a better choice.

```
| ██████████ ██████████ ██████████ |
# plot the relation between runtime and popularity
x = df['runtime']
y = df['popularity']

plt.scatter(x,y)

plt.title('Average Popularity by Different Runtime',fontsize=12)
plt.xlabel('runtime',fontsize=10)
plt.ylabel('popularity',fontsize=10)
```

```
Text(0, 0.5, 'popularity')
```



Conclusion Q2:

Combine two plots above, we can not simply say , the longer runtime, the more popular the movies are.

If the movies are within 200 minutes,it will be more popular. Once the movies run over 200 minutes, it's hard for them to gain high popularity

Q3 : Higher popularity means higher profits ?

```
# we need to get the mean of popularity
```

```
m_popularity = df['popularity'].median()
lower_popularity = df.query('popularity < {}'.format(m_popularity))
higher_popularity = df.query('popularity >= {}'.format(m_popularity))
```

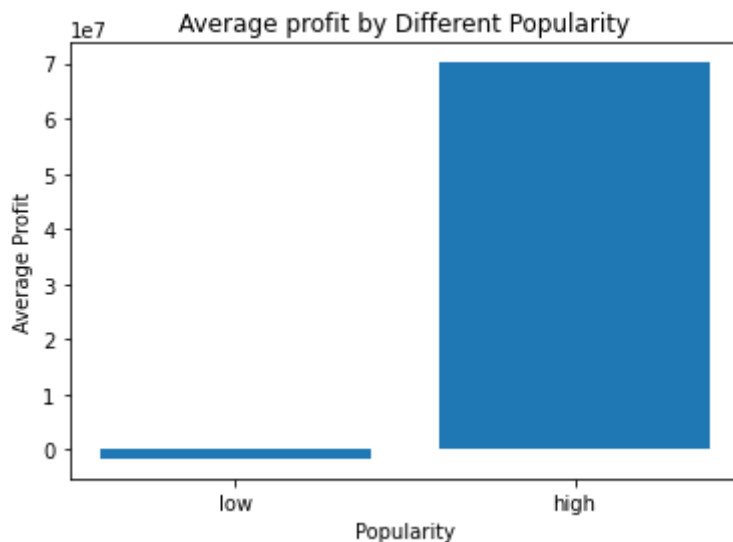
Double-click (or enter) to edit

```
# create a new column called profit. profit = Revenue - budget
df['profit'] = df['revenue'] - df['budget']
#df['profit'].head(20)
#df.head()

# average net profit for low_popularity and high_popularity
mean_profit_of_low_popularity = lower_popularity['profit'].mean()
mean_profit_of_high_popularity = higher_popularity['profit'].mean()
# df.head()

# create a bar chart with the values we get above
locations = [1,2]
heights = [mean_profit_of_low_popularity, mean_profit_of_high_popularity]
labels=['low','high']
plt.bar(locations, heights, tick_label = labels)
plt.title('Average profit by Different Popularity')
plt.xlabel('Popularity')
plt.ylabel('Average Profit')
```

```
Text(0, 0.5, 'Average Profit')
```



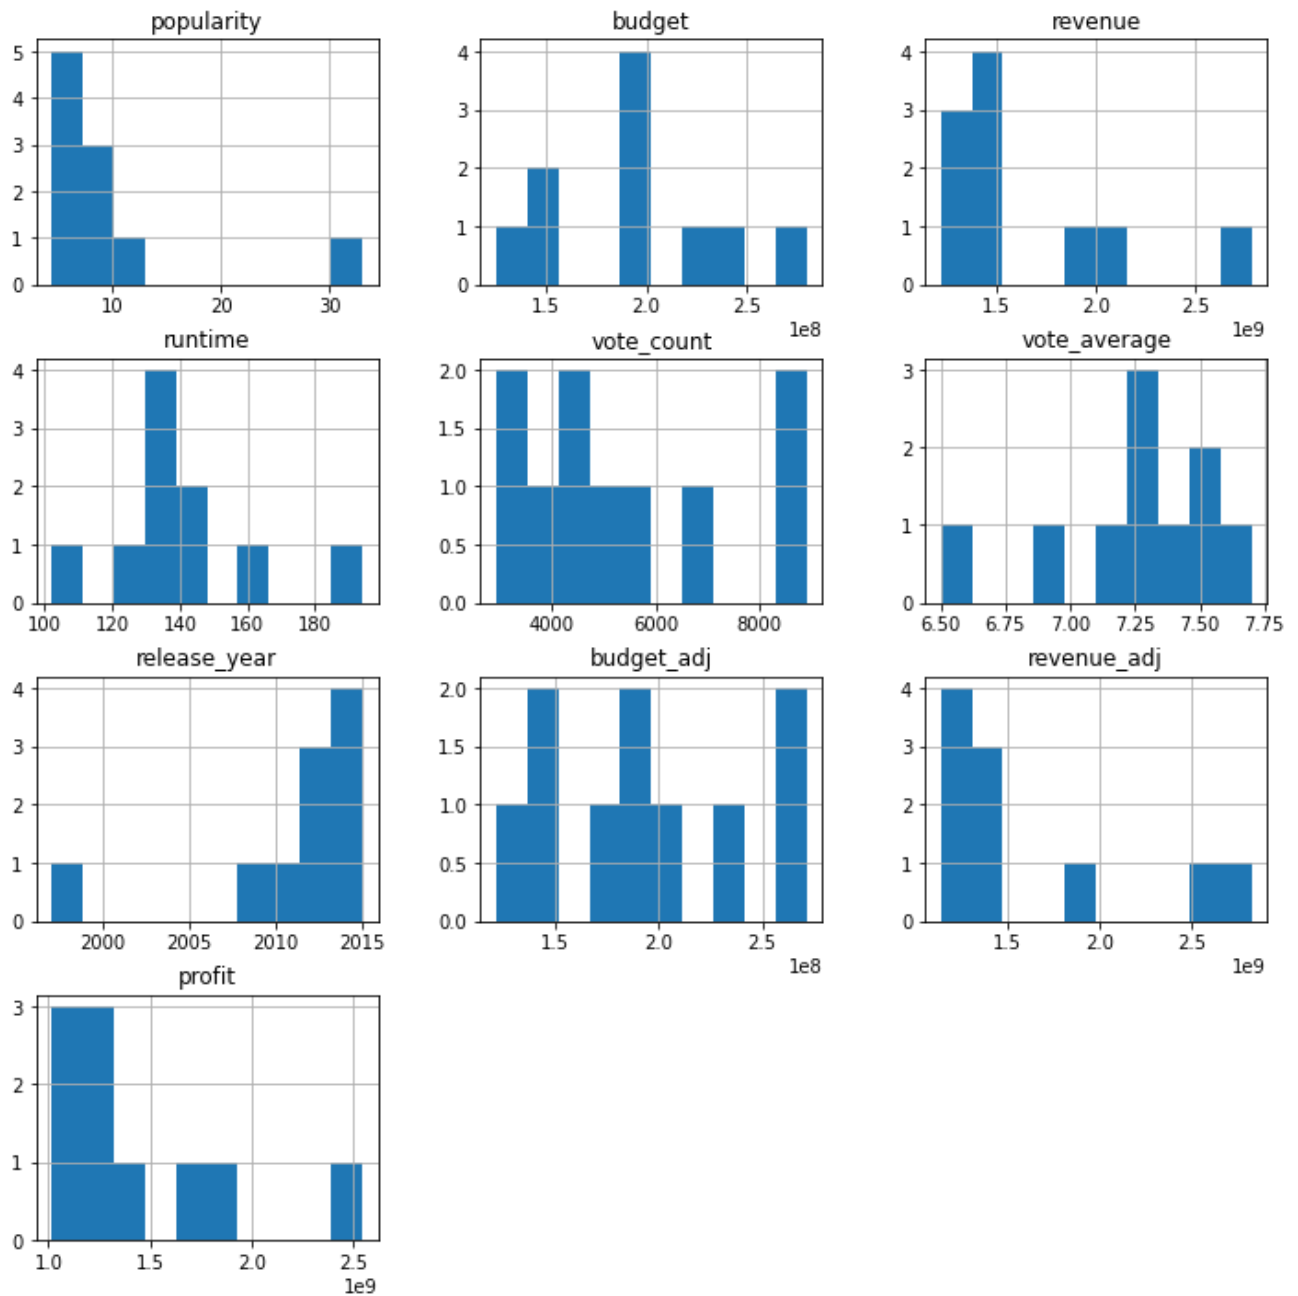
Conclusion for Question 3: as we can see above, higher popularity does make much higher average profits.

4. What Features are Associate with Top 10 Revenue Movies ?

```
top10_revenue = df.nlargest(10, 'revenue')
```

```
top10_revenue.hist(figsize=(12,12))
```

```
array([[<matplotlib.axes._subplots.AxesSubplot object at 0x7f1aae1811d0>,  
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f1aad258cd0>,  
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f1aad2253d0>],  
      [<matplotlib.axes._subplots.AxesSubplot object at 0x7f1aad1dc790>,  
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f1aad191bd0>,  
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f1aad156110>],  
      [<matplotlib.axes._subplots.AxesSubplot object at 0x7f1aad10c650>,  
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f1aad0c3a90>,  
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f1aad0c3ad0>],  
      [<matplotlib.axes._subplots.AxesSubplot object at 0x7f1aad088110>,  
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f1aad075a10>,  
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f1abd5bde90>]],  
      dtype=object)
```



Double-click (or enter) to edit

Conclusion for question 4:

There are some characteristics we can conclude from the top 10 movies. Runtime ranges from 100 mins to 200 mins. The released year are between 1995 to 2015.

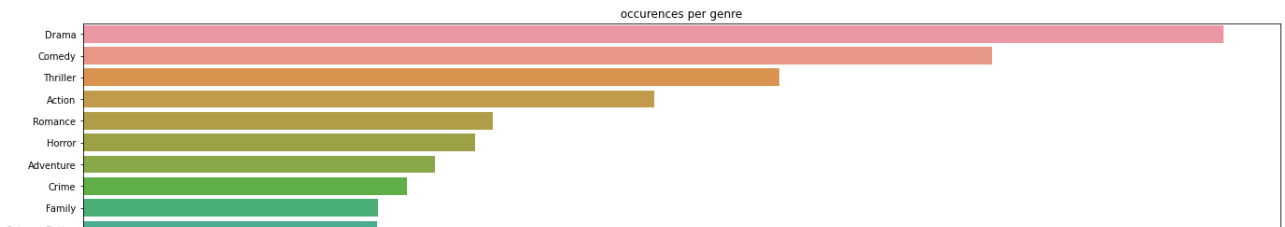
5. Which genres are most popular from year to year?

```
#The following function can give all the counts for per category
def extract_data(column_name):
    data = df[column_name].str.cat(sep = '|') # put all the genres into a long str
    # Create pandas series and store the values separately
    data = pd.Series(data.split('|')) # split the genres by |
    # Display value count in descending order
    count = data.value_counts(ascending = False) # count the occurrence of each ge
    return count

# use the function created above to split genres and count the occurrence of each
genre_count = extract_data('genres')

#create a separate dataframe to plot
df_genre_counts = pd.DataFrame({'genres': genre_count.index, 'count': genre_count.}
#df_genre_counts

f, ax = plt.subplots(figsize=(23, 9))
# use the dataframe just created as the input data
sns.barplot(x = 'count', y = 'genres', data=df_genre_counts) # how to get the data
ax.set_title(' occurrences per genre ')
ax.set_xlabel('occurrences')
ax.set_ylabel('genres')
plt.show()
```



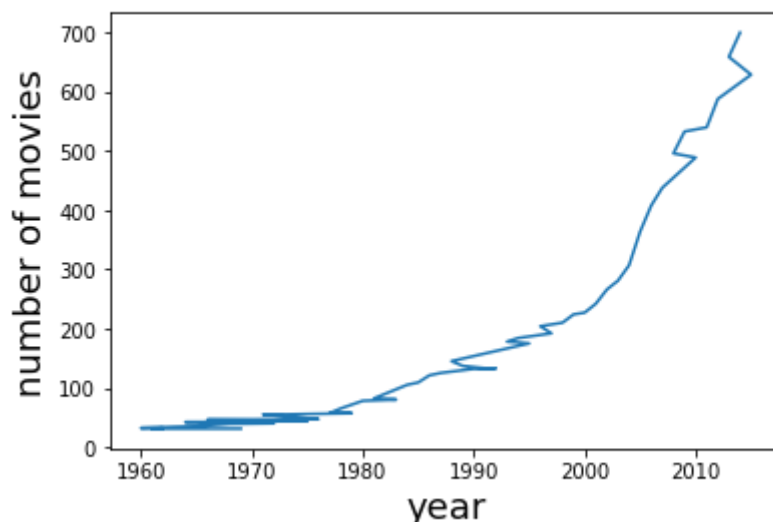
```
director_count = extract_data('director')
director_count
```

```
Woody Allen      46
missing          44
Clint Eastwood   34
Martin Scorsese  31
Steven Spielberg 30
..
Mike Maguire     1
Tom Kuntz        1
John Simpson     1
Simon Hunter     1
Harold P. Warren 1
Length: 5363, dtype: int64
```

```
movie_count = df['release_year'].value_counts()
# movie_count.plot(xlabel='year',ylabel='number of movies',title='Number of Movies
fig = plt.figure()
plt.plot(movie_count)
fig.suptitle('Number of Movies Released Each Year',fontsize=20)
plt.xlabel('year',fontsize=18)
plt.ylabel('number of movies',fontsize=18)
```

```
Text(0, 0.5, 'number of movies')
```

Number of Movies Released Each Year



Through above two plots, we can see The top 5 genres are Drama, Comedy, Action, Horror and Adventure. The number of movies increased along the time.

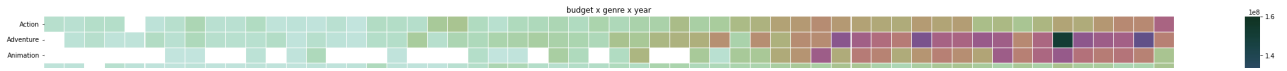
```
# The following is a really comprehensive plot. It shows the revenue and budget for
# genres are so specific, I will just take the first genre for each movie instead of
```

```
df['genre'] = df['genres'].apply(lambda x: x.split('|')[0])

# plot all the genre types for each year with the budget and revenue
genre_year = df.groupby(['genre', 'release_year']).mean().sort_index()
df_gyBudget = genre_year.pivot_table(index=['genre'], columns=['release_year'], va
df_gyBudget = genre_year.pivot_table(index=['genre'], columns=['release_year'], va

df_gyGross = genre_year.pivot_table(index=['genre'], columns=['release_year'], val
f, [axA, axB] = plt.subplots(figsize=(40, 20), nrows=2)
cmap = sns.cubehelix_palette(start=1.5, rot=1.5, as_cmap=True)
sns.heatmap(df_gyBudget, xticklabels=3, cmap=cmap, linewidths=0.05, ax=axA)
sns.heatmap(df_gyGross, xticklabels=3, cmap=cmap, linewidths=0.05, ax=axB)
axA.set_title('budget x genre x year')
axA.set_xlabel('release_years')
axA.set_ylabel('genres')

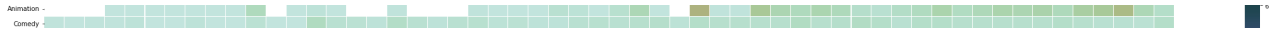
axB.set_title('revenue x genre x year')
axB.set_xlabel('release_years')
axB.set_ylabel('genres')
plt.show()
```



Conclusion for Question 5: As the time goes, we have a wider range of movies and genres to choose from. From 1984 to 2014, there are more and more high budget, high revenue movies. But compared to the budgets,



▼ Conclusion:



Based on the analysis I did above, we can make the following summarizations:

1. The quantity and range of movie gets larger. We have more choices to choose from as an audience.
2. We can not say high budget guarantees high popularity. But for movies with higher budgets do produce higher average popularity.
3. To produce a more popular movie, the runtime should be best around 150 mins; Drama, Comedy, Action, these genres would be preferable.

▼ Limitations:

1. These are factors that makes the movies become popular and successful. But we should also notice the limitations. There are some missing data and many erroneous zeros which may affect the analysis.
2. It's hard for us to know how the vote_counts and popularity are measured.
3. For foreign movies, currency is not indicated. inflation over the years should also be taken into consideration.

Reference:

1. <https://pandas.pydata.org/pandas-docs/version/0.17.0/generated/pandas.DataFrame.nlargest.html>
2. <https://www.kaggle.com/diegoinacio/imdb-genre-based-analysis>
3. <https://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.apply.html>
4. <https://pandas.pydata.org/pandas-docs/stable/visualization.html>

✓ 0s completed at 10:22 AM ● ✕